

ISSR Short Course

Gregory J. Matthews ¹

¹Department of Mathematics and Statistics
Loyola University Chicago

June 2015

Outline

Web scraping

More web scraping

Reading data from URL

R can read data directly from the internet if the URL provides links directly to a data source that can be read by R.

```
#Used to be able to do this:  
#titanic<-read.csv("http://bit.ly/1KSjVqg")  
#Now you have to do this:  
library(repmis)  
  
## Warning:  package 'repmis' was built under R  
version 3.1.3  
  
titanic<-source_data("http://bit.ly/1KSjVqg")  
  
## Downloading data from:  http://bit.ly/1KSjVqg  
##  
## SHA-1 hash of the downloaded data file is:  
## ea08b483790c2a7bc9b95b0f923526f8e60eae44  
  
class(titanic)
```

R XML package

- ▶ Usually this isn't the case though.
- ▶ Let's say we want to read data from a table that is in a web page.
- ▶ We can use functions in the R package XML to help here.

R XML package

Simple Example:

```
library(XML)
#Let's go look at this page
url<-"http://www.baseball-reference.com/teams/BOS/2014.shtml"
#Convert the web page tables into R data.frames
redSoxTable<-readHTMLTable(url)
class(redSoxTable)

## [1] "list"

names(redSoxTable)

## [1] "team_batting"          "team_pitching"
## [3] "standard_fielding"     "players_value_batting"
## [5] "players_value_pitching"

class(redSoxTable$team_batting)

## [1] "data.frame"
```

- ▶ The object `redSoxTable` is of class `list`.
- ▶ Each element of the list contains a `data.frame` object that corresponds to one of the tables on the web page.

R XML package

```
#Let's look in
```

```
redSoxTable$team_batting[1:3,1:8]
```

##	Rk	Pos	Name	Age	G	PA	AB	R
## 1	1	C	A.J. Pierzynski*	37	72	274	256	19
## 2	2	1B	Mike Napoli	32	119	500	415	49
## 3	3	2B	Dustin Pedroia	30	135	609	551	72

- ▶ Some of the player's names have characters in them that I don't want. How do I get rid of '*' and '#' in some player's names?
- ▶ What if I only want to look at infielders?

R XML package

```
#Use gsub to remove unwanted character
redSoxTable$team_batting[,3]<-
  gsub("[*#]", "", redSoxTable$team_batting[,3])
#Unwanted characters are removed
redSoxTable$team_batting[1:3,1:8]
```

##	Rk	Pos	Name	Age	G	PA	AB	R
## 1	1	C	A.J. Pierzynski	37	72	274	256	19
## 2	2	1B	Mike Napoli	32	119	500	415	49
## 3	3	2B	Dustin Pedroia	30	135	609	551	72

R XML package

```
#To get just infielders, we use the function %in%  
#Only pull out infielders  
infInd <- redSoxTable$team_batting$Pos%in%c("1B", "2B", "3B")  
infielders <- redSoxTable$team_batting[infInd,]  
infielders[,1:8]
```

##	Rk	Pos	Name	Age	G	PA	AB	R
## 2	2	1B	Mike Napoli	32	119	500	415	49
## 3	3	2B	Dustin Pedroia	30	135	609	551	72
## 4	4	SS	Xander Bogaerts	21	144	594	538	60
## 5	5	3B	Will Middlebrooks	25	63	234	215	14
## 17	16	SS	Stephen Drew	31	39	145	131	11
## 27	26	3B	Ryan Roberts	33	8	22	19	1
## 29	28	1B	Ryan Lavarney	26	9	10	10	0

R XML package

- ▶ Now let's say I want to do something across every team.
- ▶ How do I do that?
- ▶ Notice that in the URL the only piece that changes is the team abbreviation.
- ▶ [http://www.baseball-reference.com/teams/**BOS**/2014.shtml](http://www.baseball-reference.com/teams/BOS/2014.shtml)
- ▶ So first, let's get a vector of these abbreviations.
- ▶ We can do this by going to this url:
<http://www.baseball-reference.com/leagues/MLB/2014.shtml>

R XML package

```
teamsData<-readHTMLTable("http://www.baseball-reference.com/leagues/MLB/2014.shtml")
class(teamsData)

## [1] "list"

names(teamsData)

## [1] "NULL"                "teams_standard_batting"
## [3] "teams_standard_pitching" "team_output"
## [5] "teams_standard_fielding"

teamsData[["teams_standard_batting"]][1:3,1:5]

##      Tm #Bat BatAge R/G   G
## 1 ARI   52   27.6 3.80 162
## 2 ATL   39   26.8 3.54 162
## 3 BAL   44   28.3 4.35 162

#Construct a vector of all the team abbreviations
teams<-as.character(teamsData[["teams_standard_batting"]]$Tm[1:30])
teams

## [1] "ARI" "ATL" "BAL" "BOS" "CHC" "CHW" "CIN" "CLE" "COL" "DET" "HOU"
## [12] "KCR" "LAA" "LAD" "MIA" "MIL" "MIN" "NYM" "NYY" "OAK" "PHI" "PIT"
## [23] "SDP" "SEA" "SFG" "STL" "TBR" "TEX" "TOR" "WSN"
```

R XML package

```
#resList will store the results
resList<-list()
#We will loop over all teams
for (t in teams){print(t)
  url<-paste("http://www.baseball-reference.com/teams/",t,"/2014.shtml",sep="")
  teamData<-readHTMLTable(url)
  #Pull out the team batting component
  output<-teamData$team_batting
  #Convert player name to a character
  output[,3]<-as.character(output[,3])
  output[,3]<-gsub("[*#]", "", output[,3])
  #Store the output
  resList[[t]]<-output[output[,3]!="",]
}
```

R XML package

```
class(resList)
names(resList)
resList$BOS[1:3,1:8]
resList$NYY[1:3,1:8]
```

R XML package

```
#Now I want to stack all of these  
#data frames on top of each other  
allTeams<-do.call(rbind,resList)  
dim(allTeams)
```

- ▶ So far we have scraped data from tables online.
- ▶ What if we want to scrap unstructured data from the web.

- ▶ This first example involves scraping the content of presidential inaugural speeches.
- ▶ We're going to use the functions `getURL` and `parseHTML` here.


```
library(RCurl)
```

```
## Loading required package: bitops
```

```
library(XML)
```

```
presList<-list()
```

```
url<-"http://www.presidentialrhetoric.com/historicspeeches/bush/
```

- ▶ We can read the content of the webpage.
- ▶ But it is a mess right now.

```
a<-getURL(url)
str(a)
```

- ▶ If we want to parse the html we can use the following commands.
- ▶ Either way (htmlParse or getURL) we can extract the text we want.

```
b<-htmlParse(a)
b
```

Before we were reading HTML tables from the web. Now we are reading an HTML table that is an object in R.

```
x<-readHTMLTable(a)
```

```
x
```

let's clean it up.

```
#clean it up a bit
text<-levels(x[[6]]$V1)
text<-gsub("\n","",text)
text<-gsub("[,\\.]", "",text)
#Make it a plain text document and add it to a list
library(tm)
presList[["HWBush"]]<-PlainTextDocument(text)
```

- ▶ Now we have one speech in a list.
- ▶ Let's go get some more speeches and make a corpus.
- ▶ We're going to repeat what we just did for Clinton and W. Bush.
- ▶ Then we will make these three speeches into a Corpus.

```
url<-"http://www.presidentialrhetoric.com/historicspeeches/  
a<-getURL(url)  
b<-htmlParse(url)  
x<-readHTMLTable(b)  
text<-levels(x[[6]]$V1)  
text<-gsub("\n","",text)  
text<-gsub("[,.]","",text)  
presList[["Clinton"]]<-PlainTextDocument(text)
```

```
url<-"http://www.presidentialrhetoric.com/historicspeeches/  
a<-getURL(url)  
b<-htmlParse(url)  
x<-readHTMLTable(b)  
text<-levels(x[[6]]$V1)  
text<-gsub("\n", "", text)  
text<-gsub("[,.] ", "", text)  
presList[["Bush"]]<-PlainTextDocument(text)
```


- ▶ We have a list object called `presList`.
- ▶ Each element of the list is a `PlainTextDocument`.
- ▶ We can now turn the list into a `Corpus`.

```
presList<-Corpus(VectorSource(presList))  
class(presList)
```

Basic Filtering

- ▶ Remove extra white space
- ▶ Remove “stop words” (i.e. “a”, “an”, and “the”)
- ▶ Stemming: Convert words to their stems (i.e “Fishing” and “Fished” become “Fish”)
- ▶ Convert to lower case: This way “Dog” at the beginning of a sentence is treated the same was as “dog” in the middle of the sentence.
- ▶ Remove sparse terms: remove terms that are only used in a small number of documents.

Transformations on Corpora

```
#Removes extra whitespace
presList<-tm_map(presList,stripWhitespace)
#Stop words: Words filtered out before processing.
stopwords("english")[1:10]
presList<-tm_map(presList, removeWords, stopwords("english"))
#presList <- tm_map(presList, tolower) #OLD
presList <- tm_map(presList, content_transformer(tolower)) #NEW
#Stemming: reducing words to their stems.
presList<-tm_map(presList, stemDocument) #requires SnowballC package
presList<-tm_map(presList, PlainTextDocument)
```

```
presTDM<-TermDocumentMatrix(presList)
presTDM<-removeSparseTerms(presTDM,0.5)
```

#1 means it has to be in all documents
#0 means we keep all words
#0.5 means we keep words that appear in
#at least half of the documents.

Term Frequency

- ▶ We may be interested in seeing how often certain words appear.
- ▶ This is analagous to the table function in R, but we can filter out certain terms that we may not want based on some criteria.

```
#Set control parameters
ctrl<-list(removePunctuation =
            list(preserve_intra_word_dashes = TRUE),
            stopwords=c("bush", "Bush"),
            stemming=TRUE,
            wordLengths=c(4, Inf))
#Look at a frequency table of words
sort(termFreq(presList[["HWBush"]]), control=ctrl)
```

#We can also do this:

```
findFreqTerms(presList[["HWBush"]],lowfreq=3,highfreq=10)
```

- ▶ We can also measure the similarity or dissimilarity of the texts in the corpus.
- ▶ We also may be interested in looking for associations of words within certain texts. Simply how often do these words appear in the same text.


```
library(proxy)
#Slot 6 contains the text.
presTDM<-TermDocumentMatrix(presList)
#
dissimilarity(presTDM,method="eJaccard")
findAssocs(presTDM,c("responsibility","vision"),corlimit=c
```

- ▶ Right now we have three presidential inaugural addresses.
- ▶ But what if we want to get ALL the presidential inaugural addresses?
- ▶ We need to provide URLs for each speech and this can be tedious.
- ▶ Luckily, we can use R to get the list of URLs for us.

```
#Now get them all  
#pull in the content of the page  
#the parse the HTML  
doc<-htmlParse(getURL("http://www.presidentialrhetoric.com/  
#This next line pulls out all of the links  
vec<-xpathSApply(doc, "//a/@href")  
#Now we can pull out only the speeches that were inaugural  
vecList<-vec[grep("inaugural",vec)]
```

- ▶ What do we have now?
- ▶ We have a vector containing all of the last pieces of the URLs that we will need for each presidential inaugural address.
- ▶ We can combine these pieces with the root URL using paste and loop across all of the speeches to get all of them.

```
presList<-list()
for (vvv in vecList){print(vvv)
url<-paste("http://www.presidentialrhetoric.com/historicspe
a<-getURL(url)
b<-htmlParse(url)
x<-readHTMLTable(b)
text<-levels(x[[6]]$V1)
text<-gsub("\n","",text)
text<-gsub("[,.]","",text)
presList[[vvv]]<-PlainTextDocument(text)
}
```

```
presCorpus <- Corpus(VectorSource(presList))
presCorpus <- tm_map(presCorpus, content_transformer(tolower))
presCorpus<-tm_map(presCorpus,removeWords,stopwords("english"))
presCorpus<-tm_map(presCorpus,removePunctuation)
presCorpus<-tm_map(presCorpus,removeNumbers)
presCorpus<-tm_map(presCorpus,stripWhitespace)
presTDM<-TermDocumentMatrix(presCorpus)
presTDM$dimnames$Docs<-names(presList)
```

```
findFreqTerms(presTDM, 500)
#[1] "government" "people"      "will"
findAssocs(presTDM, "government", 0.7)
#
#      government
#system      0.78
#states      0.73
#interests   0.72
#protection  0.71
#adoption    0.70
```

Let's try it:

- ▶ Get inaugural speeches of the last three presidents.
- ▶ Turn these speeches into a corpus.
- ▶ Get all inaugural presidential speeches.
- ▶ Turn these speeches into a corpus.
- ▶ Find frequently used terms.
- ▶ See BushCode.R